

UNIs: sg3637, um2147

Task 1

In [210]:

```
1 import pandas as pd
2 import warnings
3 from IPython.display import display
4 warnings.filterwarnings('ignore')
5 pd.options.display.max_columns = None
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8 import datetime as dt
9 import numpy as np
10 warnings.simplefilter('ignore')
```

In [4]:

```
1 #Reading Sample from the Research Data Table
2 import random
3
4 filename = "data/OP_DTL_RSRCH_PGYR2017_P01182019.csv"
5 n = sum(1 for line in open(filename)) - 1
6 #number of records in file (excludes header)
7 s = 100000
8 #desired sample size
9 skip = sorted(random.sample(range(1,n+1),n-s))
10 #the 0-indexed header will not be included in the skip list
11 dfResearch = pd.read_csv(filename, skiprows=skip)
12
13 dfResearch.shape
```

Out[4]:

(100000, 176)

In [5]:

```
1 dfResearch.head()
```

Out[5]:

Recipient_Primary_Business_Street_Address_Line2	Recipient_City	Recipient_State	Recipient_Zip
NaN	Fridley	MN	
Suite 255	San Diego	CA	
NaN	KALAMAZOO	MI	
NaN	LAWRENCEVILLE	NJ	
NaN	JEFFERSONVILLE	IN	

In [6]:

```
1 #Viewing all the Columns from dfResearch
2 dfResearch.columns
```

Out[6]:

```
Index(['Change_Type', 'Covered_Recipient_Type',
      'Noncovered_Recipient_Entity_Name', 'Teaching_Hospital_CCN',
      'Teaching_Hospital_ID', 'Teaching_Hospital_Name',
      'Physician_Profile_ID', 'Physician_First_Name', 'Physician_Middle_Name',
      'Physician_Last_Name',
      ...,
      'Preclinical_Research_Indicator', 'Delay_in_Publication_Indicator',
      'Name_of_Study', 'Dispute_Status_for_Publication', 'Record_ID',
      'Program_Year', 'Payment_Publication_Date',
      'ClinicalTrials_Gov_Identifier', 'Research_Information_Link',
      'Context_of_Research'],
      dtype='object', length=176)
```

In [7]:

```
1 #As per the data dictionary, "Physicians may be identified as covered recipients
2 # with research-related payment records. Hence, we will replace the missing val
3
4 dfResearch['Physician_Primary_Type'].fillna(dfResearch['Principal_Investigator_1
5 dfResearch['Physician_Specialty'].fillna(dfResearch['Principal_Investigator_1_Sp
6 dfResearch['Physician_License_State_code1'].fillna(dfResearch['Principal_Investi
```

In [8]:

```
1 #Read from Chunk with condition
2 #https://stackoverflow.com/questions/34549402/pandas-read-a-small-random-sample
```

In [9]:

```
1 #Reading Sample for the other table
2 import random
3
4 filename = "data/OP_DTL_GNRL_PGYR2017_P01182019.csv"
5 n = sum(1 for line in open(filename)) - 1
6 #number of records in file (excludes header)
7 s = 100000
8 #desired sample size
9 skip = sorted(random.sample(range(1,n+1),n-s))
10 #the 0-indexed header will not be included in the skip list
11 df = pd.read_csv(filename, skiprows=skip)
```

In [10]:

```
1 #df consists of payments made in general/ not via research. Adding the label: I.
2 df['Label'] = 0
3 dfResearch['Label'] = 1
```

In [34]:

```
1 #Combining the two dataframes in df_final
2 df_final = pd.concat([dfResearch, df], join="inner")
3 df_final = df_final.sample(frac=1, random_state=42).reset_index(drop=True)
4
5 #Dropping columns with all na values
6 df_final.dropna(axis=1, how='all', inplace = True)
```

Identifying Features

In the initial feature selection, we will use heuristics to reduce the of features that evidently would not impact the model. We will then use more statistical methods and visulization to further select on the relevant features.

1. We will drop all the features that have greater than 50% null values.

In [35]:

```
1 #Dropping Features that have >50% null values
2 df_final = df_final.loc[:, df_final.isnull().mean() <= .5]
```

In [36]:

```
1 df_final.head()
```

Out[36]:

	Change_Type	Covered_Recipient_Type	Physician_Profile_ID	Physician_First_Name	Physician_L
0	UNCHANGED	Covered Recipient Physician	3297.0	JOHN	
1	UNCHANGED	Non-covered Recipient Entity	NaN	NaN	
2	UNCHANGED	Covered Recipient Physician	180292.0	KEVIN	
3	UNCHANGED	Non-covered Recipient Entity	NaN	NaN	
4	UNCHANGED	Non-covered Recipient Entity	NaN	NaN	

In [37]:

```
1 df_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 33 columns):
Change_Type                                                    20
0000 non-null object
Covered_Recipient_Type                                          20
0000 non-null object
Physician_Profile_ID                                           10
3722 non-null float64
Physician_First_Name                                           10
3721 non-null object
Physician_Last_Name                                            10
3721 non-null object
Recipient_Primary_Business_Street_Address_Line1              19
9879 non-null object
Recipient_City                                                  19
9879 non-null object
Recipient_State                                                 19
9799 non-null object
Recipient_Zip_Code                                              19
9799 non-null object
Recipient_Country                                              19
9879 non-null object
Physician_Primary_Type                                          19
8794 non-null object
Physician_Specialty                                             19
8526 non-null object
Physician_License_State_code1                                  19
8791 non-null object
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name     20
0000 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_ID   20
0000 non-null int64
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name 20
0000 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State 18
8624 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country 20
0000 non-null object
Related_Product_Indicator                                       20
0000 non-null object
Covered_or_Noncovered_Indicator_1                              18
1534 non-null object
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1     16
3439 non-null object
Product_Category_or_Therapeutic_Area_1                        15
8380 non-null object
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1      16
2667 non-null object
Associated_Drug_or_Biological_NDC_1                            12
1608 non-null object
Total_Amount_of_Payment_USDollars                              20
0000 non-null float64
Date_of_Payment                                                20
0000 non-null object
Form_of_Payment_or_Transfer_of_Value                          20
0000 non-null object
```

```

Delay_in_Publication_Indicator      20
0000 non-null object
Dispute_Status_for_Publication      20
0000 non-null object
Record_ID                           20
0000 non-null int64
Program_Year                        20
0000 non-null int64
Payment_Publication_Date            20
0000 non-null object
Label                               20
0000 non-null int64
dtypes: float64(2), int64(4), object(27)
memory usage: 50.4+ MB

```

2. We will remove columns that are evidently dependent on each other. For instance, physician_id will be correlated with Physician_First_Name and Physician_Last_Name. In this case, we will remove Physician_First_Name and Physician_Last_Name. Similar approach has been taken for some of the other features.

In [38]:

```

1  #Removing some obvious identifier/name columns that can lead to information leak
2  df_final.drop(columns=['Physician_First_Name', 'Physician_Last_Name', 'Record_ID'])

```

In [39]:

```

1  df_final.head()

```

Out[39]:

	Change_Type	Covered_Recipient_Type	Physician_Profile_ID	Recipient_Primary_Business_Street
0	UNCHANGED	Covered Recipient Physician	3297.0	2525 W
1	UNCHANGED	Non-covered Recipient Entity	NaN	530 NE
2	UNCHANGED	Covered Recipient Physician	180292.0	1435 N Ra
3	UNCHANGED	Non-covered Recipient Entity	NaN	12
4	UNCHANGED	Non-covered Recipient Entity	NaN	2141 E

In [40]:

```
1 df_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 30 columns):
Change_Type                                20
0000 non-null object
Covered_Recipient_Type                    20
0000 non-null object
Physician_Profile_ID                      10
3722 non-null float64
Recipient_Primary_Business_Street_Address_Line1  19
9879 non-null object
Recipient_City                            19
9879 non-null object
Recipient_State                           19
9799 non-null object
Recipient_Zip_Code                        19
9799 non-null object
Recipient_Country                         19
9879 non-null object
Physician_Primary_Type                    19
8794 non-null object
Physician_Specialty                      19
8526 non-null object
Physician_License_State_code1             19
8791 non-null object
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name  20
0000 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_ID  20
0000 non-null int64
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name  20
0000 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State  18
8624 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country  20
0000 non-null object
Related_Product_Indicator                 20
0000 non-null object
Covered_or_Noncovered_Indicator_1         18
1534 non-null object
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1  16
3439 non-null object
Product_Category_or_Therapeutic_Area_1    15
8380 non-null object
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1  16
2667 non-null object
Associated_Drug_or_Biological_NDC_1        12
1608 non-null object
Total_Amount_of_Payment_USDollars         20
0000 non-null float64
Date_of_Payment                           20
0000 non-null object
Form_of_Payment_or_Transfer_of_Value      20
0000 non-null object
Delay_in_Publication_Indicator            20
0000 non-null object
Dispute_Status_for_Publication            20
0000 non-null object
```

Program_Year	20
0000 non-null int64	
Payment_Publication_Date	20
0000 non-null object	
Label	20
0000 non-null int64	
dtypes: float64(2), int64(3), object(25)	
memory usage: 45.8+ MB	

Let us now look at the number of unique values in each of the columns. For columns with categorical variables and less than 60 unique values, we will visualize the number of data points in each of the categories through bar charts. This will allow us to examine any categories that are there in only one class. We will eliminate such records as this will allow us to minimize data leakage.

In [41]:

```
1 df_final.nunique().sort_values()
```

Out[41]:

```
Program_Year
1
Delay_in_Publication_Indicator
1
Payment_Publication_Date
1
Label
2
Dispute_Status_for_Publication
2
Covered_or_Noncovered_Indicator_1
2
Related_Product_Indicator
2
Change_Type
3
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1
4
Covered_Recipient_Type
4
Form_of_Payment_or_Transfer_of_Value
4
Physician_Primary_Type
6
Recipient_Country
7
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country
26
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State
43
Physician_License_State_code1
54
Recipient_State
56
Physician_Specialty
275
Date_of_Payment
365
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name
874
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_ID
979
Associated_Drug_or_Biological_NDC_1
993
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name
993
Product_Category_or_Therapeutic_Area_1      1
060
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1      3
941
Recipient_City      8
122
Recipient_Zip_Code      38
028
Total_Amount_of_Payment_USDollars      46
```

```
399
Recipient_Primary_Business_Street_Address_Line1          73
954
Physician_Profile_ID          73
964
dtype: int64
```

In [42]:

```
1  #Dropping the columns that have only 1 unique value = Zero Variance
2  df_final.drop(columns=['Program_Year', 'Delay_in_Publication_Indicator', 'Payme
```

In [43]:

```
1 df_final.nunique().sort_values()
```

Out[43]:

```
Label
2
Covered_or_Noncovered_Indicator_1
2
Related_Product_Indicator
2
Dispute_Status_for_Publication
2
Change_Type
3
Form_of_Payment_or_Transfer_of_Value
4
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1
4
Covered_Recipient_Type
4
Physician_Primary_Type
6
Recipient_Country
7
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country
26
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State
43
Physician_License_State_code1
54
Recipient_State
56
Physician_Specialty
275
Date_of_Payment
365
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name
874
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_ID
979
Associated_Drug_or_Biological_NDC_1
993
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name
993
Product_Category_or_Therapeutic_Area_1      1
060
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1      3
941
Recipient_City      8
122
Recipient_Zip_Code      38
028
Total_Amount_of_Payment_USDollars      46
399
Recipient_Primary_Business_Street_Address_Line1      73
954
Physician_Profile_ID      73
964
dtype: int64
```

In [44]:

```
1 df_final.head()
```

Out[44]:

	Change_Type	Covered_Recipient_Type	Physician_Profile_ID	Recipient_Primary_Business_Street
0	UNCHANGED	Covered Recipient Physician	3297.0	2525 W
1	UNCHANGED	Non-covered Recipient Entity	NaN	530 NE
2	UNCHANGED	Covered Recipient Physician	180292.0	1435 N Ra
3	UNCHANGED	Non-covered Recipient Entity	NaN	12
4	UNCHANGED	Non-covered Recipient Entity	NaN	2141 E

In [45]:

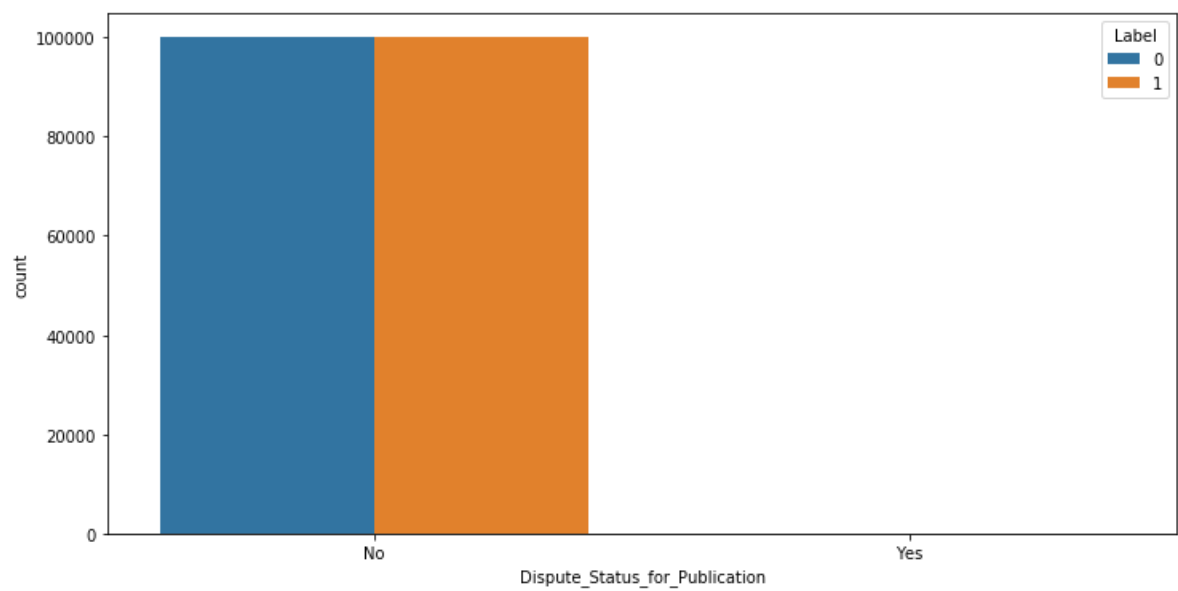
```
1 #Dispute_Status_for_Publication
2 df_final.groupby(['Label', 'Dispute_Status_for_Publication']).count()
```

Out[45]:

		Change_Type	Covered_Recipient_Type	Physician_Profile_ID
Label	Dispute_Status_for_Publication			
0	No	99989	99989	99610
	Yes	11	11	8
1	No	99955	99955	4097
	Yes	45	45	7

In [46]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Dispute_Status_for_Publication", data=df_final, hue='Label')
```



We will keep this column as it does not lead to information leakage. Both classes have 'Yes' and 'No'.

In [47]:

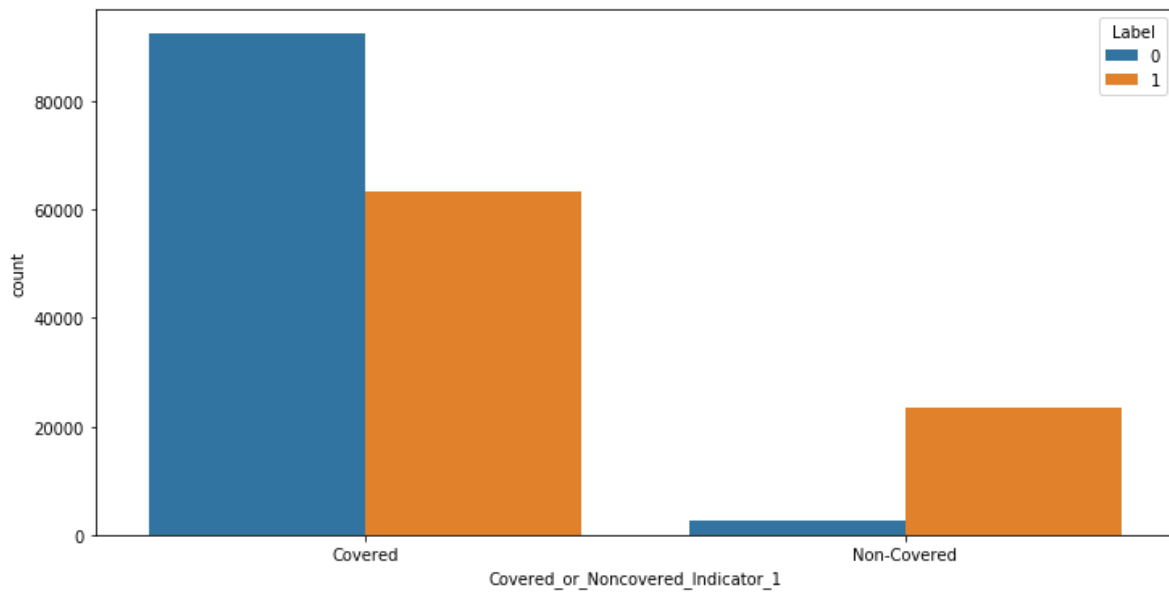
```
1 #Covered_or_Noncovered_Indicator_1
2 df_final.groupby(['Label', 'Covered_or_Noncovered_Indicator_1']).count()
```

Out[47]:

		Change_Type	Covered_Recipient_Type	Physician_Prof
Label	Covered_or_Noncovered_Indicator_1			
0	Covered	92343	92343	9
	Non-Covered	2648	2648	0
1	Covered	63205	63205	0
	Non-Covered	23338	23338	0

In [48]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Covered_or_Noncovered_Indicator_1", data=df_final, hue='Label')
```



We will keep this column as it does not lead to information leakage. Both classes have 'Covered' and 'Non-Covered'

In [49]:

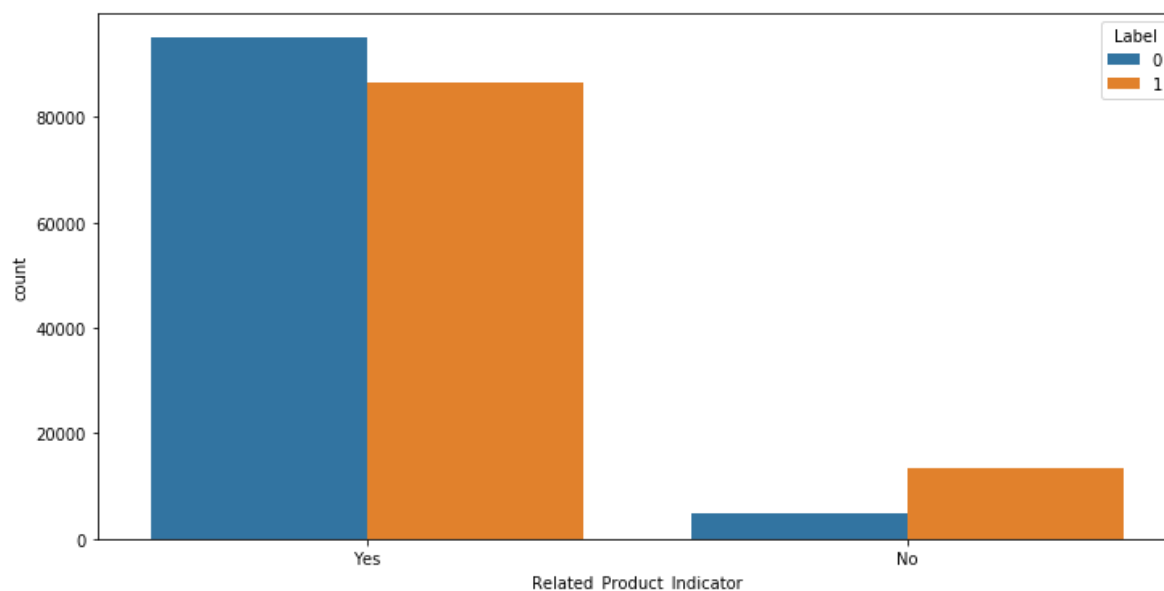
```
1 #Related_Product_Indicator
2 df_final.groupby(['Label', 'Related_Product_Indicator']).count()
```

Out[49]:

		Change_Type	Covered_Recipient_Type	Physician_Profile_ID	Re
Label	Related_Product_Indicator				
0	No	5009	5009	4915	
	Yes	94991	94991	94703	
1	No	13457	13457	503	
	Yes	86543	86543	3601	

In [50]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Related_Product_Indicator", data=df_final, hue='Label')
```



We will keep this column as it does not lead to information leakage. Both classes have 'Yes' and 'No'

In [51]:

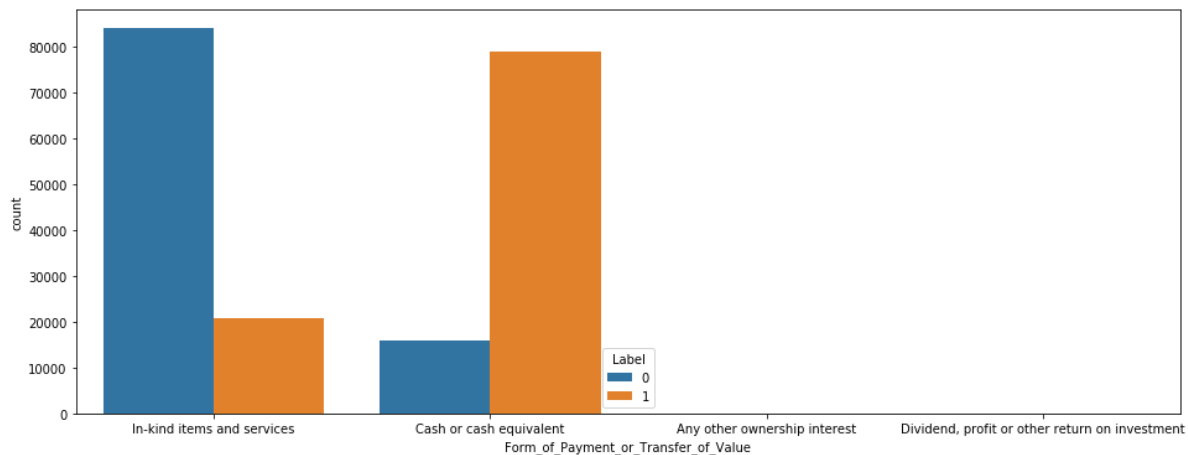
```
1 #Form_of_Payment_or_Transfer_of_Value
2 df_final.groupby(['Label', 'Form_of_Payment_or_Transfer_of_Value']).count()
```

Out[51]:

		Change_Type	Covered_Recipient_Type	Physician_P
Label	Form_of_Payment_or_Transfer_of_Value			
0	Any other ownership interest	1	1	
	Cash or cash equivalent	15882	15882	
	Dividend, profit or other return on investment	4	4	
	In-kind items and services	84113	84113	
1	Cash or cash equivalent	79129	79129	
	In-kind items and services	20871	20871	

In [52]:

```
1 plt.figure(figsize = (16,6))
2 _ = sns.countplot(x="Form_of_Payment_or_Transfer_of_Value", data=df_final, hue=
```



From the above chart, it looks like 'Form_of_Payment_or_Transfer_of_Value' is very much correlated with the target label. As such, we will drop this feature completely.

In [53]:

```
1 df_final.drop(columns='Form_of_Payment_or_Transfer_of_Value', inplace=True)
```

In [54]:

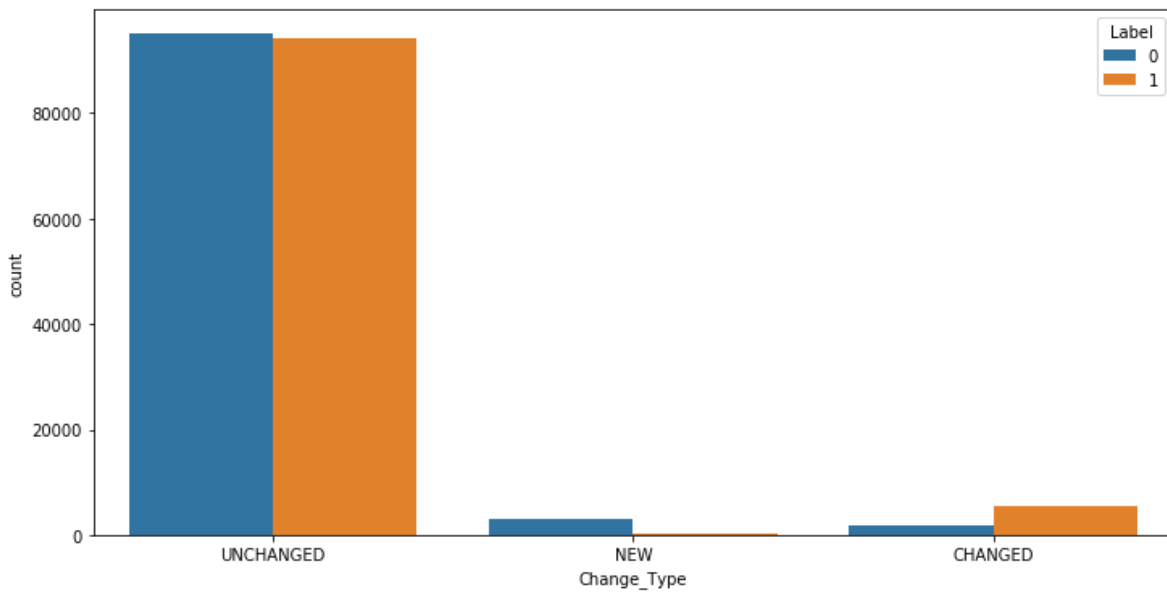
```
1 #Change_Type
2 df_final.groupby(['Label', 'Change_Type']).count()
```

Out[54]:

		Covered_Recipient_Type	Physician_Profile_ID	Recipient_Primary_Business_Si
Label	Change_Type			
0	CHANGED	1904	1899	
	NEW	3059	3048	
	UNCHANGED	95037	94671	
1	CHANGED	5470	86	
	NEW	276	123	
	UNCHANGED	94254	3895	

In [55]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Change_Type", data=df_final, hue='Label')
```



We will keep this column as it does not lead to information leakage. Both classes have same categories.

In [56]:

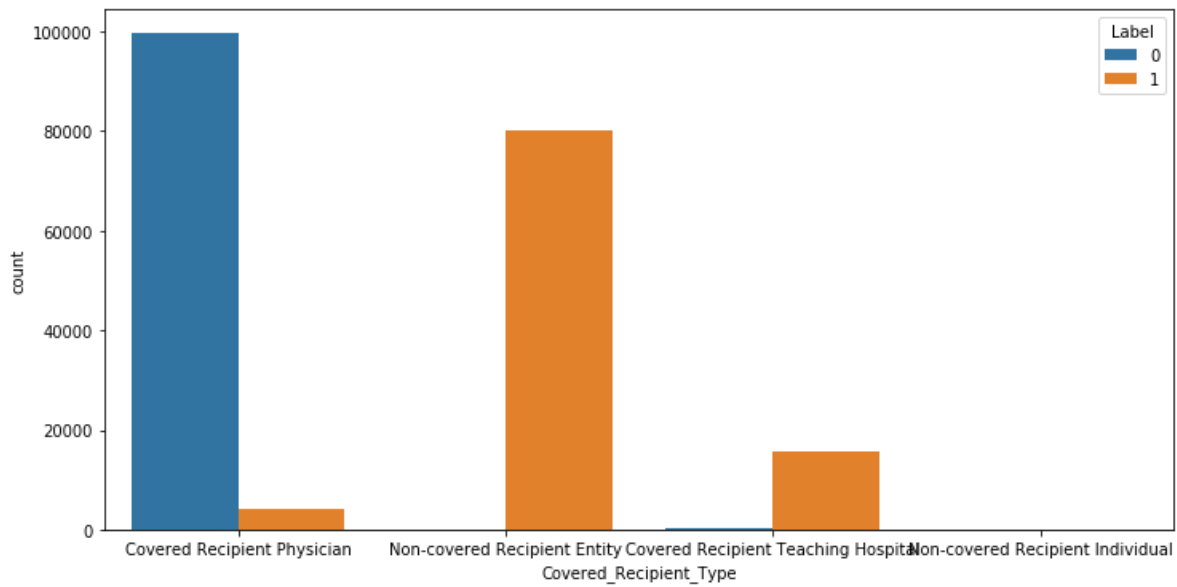
```
1 #Covered_Recipient_Type
2 df_final.groupby(['Label', 'Covered_Recipient_Type']).count()
```

Out[56]:

		Change_Type	Physician_Profile_ID	Recipient_Primary_Business_St
Label	Covered_Recipient_Type			
0	Covered Recipient Physician	99618	99618	
	Covered Recipient Teaching Hospital	382	0	
1	Covered Recipient Physician	4104	4104	
	Covered Recipient Teaching Hospital	15771	0	
	Non-covered Recipient Entity	80004	0	
	Non-covered Recipient Individual	121	0	

In [57]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Covered_Recipient_Type", data=df_final, hue='Label')
```



From the above chart, it looks like 'Covered_Recipient_Type' is very much correlated with the target label. As such, we will drop this feature completely.

In [58]:

```
1 df_final.drop(columns='Covered_Recipient_Type', inplace=True)
```

In [59]:


```
1 #Recipient_Country
2 df_final.groupby(['Label', 'Recipient_Country']).count()
```

Out[59]:

		Change_Type	Physician_Profile_ID	Recipient_Primary_Business_Street_A
Label	Recipient_Country			
0	Australia	1	1	
	Great Britain (UK)	1	1	
	United States	99997	99615	
	United States Minor Outlying Islands	1	1	
1	Belgium	1	0	
	Canada	10	0	
	Great Britain (UK)	64	0	
	Japan	1	0	
	United States	99802	4104	
	United States Minor Outlying Islands	1	0	

In [60]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Recipient_Country", data=df_final, hue='Label')
```



From the above, we can see that most of the data has Recipient_Country as United States. As such, there is very little variance. In order to avoid data leakage, we will remove this feature all together.

In [61]:

```
1 #Dropping Recipient_Country Column
2 df_final.drop(columns='Recipient_Country', inplace=True)
```

In [62]:

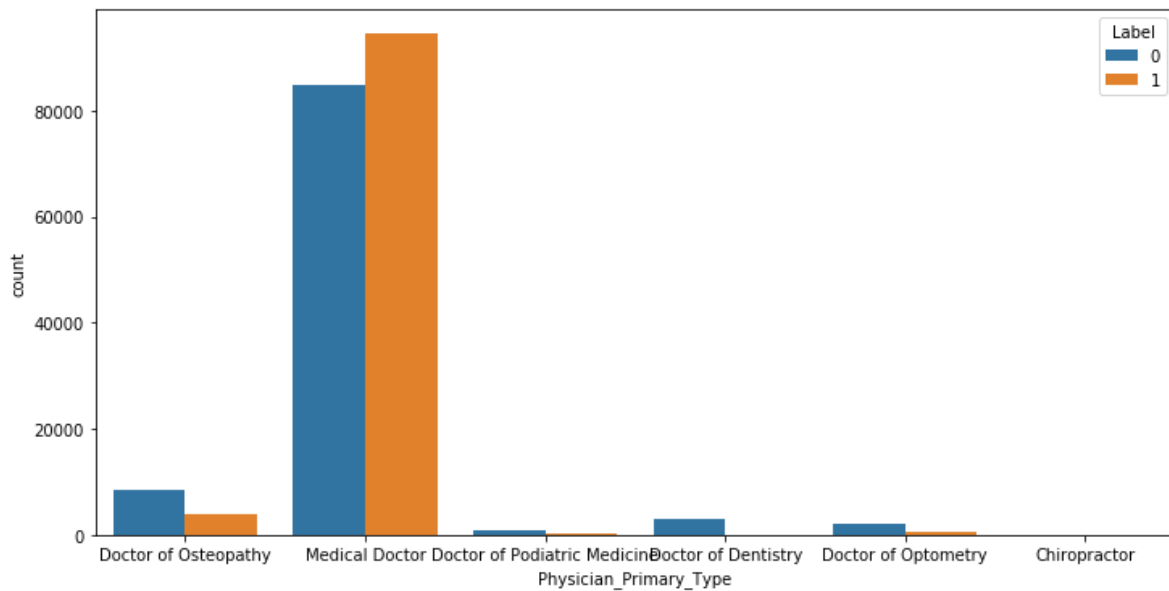
```
1 #Physician_Primary_Type
2 df_final.groupby(['Label', 'Physician_Primary_Type']).count()
```

Out[62]:

		Change_Type	Physician_Profile_ID	Recipient_Primary_Business_Sti
Label	Physician_Primary_Type			
0	Chiropractor	41	41	
	Doctor of Dentistry	3110	3110	
	Doctor of Optometry	2200	2200	
	Doctor of Osteopathy	8535	8535	
	Doctor of Podiatric Medicine	916	916	
	Medical Doctor	84816	84816	
1	Chiropractor	1	1	
	Doctor of Dentistry	98	65	
	Doctor of Optometry	459	69	
	Doctor of Osteopathy	3977	145	
	Doctor of Podiatric Medicine	219	8	
	Medical Doctor	94422	3816	

In [63]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Physician_Primary_Type", data=df_final, hue='Label')
3 #Remove this feature due to information leakage
```



From the above we can see that the Chiropractor Physician_Primary_Type is specific to label 0, hence to avoid information leakage we will remove rows where Physician_Primary_Type = Chiropractor

In [64]:

```
1 df_final = df_final[(df_final.Physician_Primary_Type != 'Chiropractor')]
```

In [65]:


```
1 #Change_Type
2 df_final.groupby(['Label', 'Change_Type']).count()
```

Out[65]:

		Physician_Profile_ID	Recipient_Primary_Business_Street_Address_Line1	Reci
Label	Change_Type			
0	CHANGED	1899		1904
	NEW	3048		3059
	UNCHANGED	94630		94996
1	CHANGED	86		5470
	NEW	123		276
	UNCHANGED	3894		94132

In [66]:

```
1 plt.figure(figsize = (12,6))
2 _ = sns.countplot(x="Change_Type", data=df_final, hue='Label')
```



We will keep this column as it does not lead to information leakage. Both classes have same categories.

In [67]:


```
1 #Recipient_State
2 df_final.groupby(['Label', 'Recipient_State']).count()
```

Out[67]:

		Change_Type	Physician_Profile_ID	Recipient_Primary_Business_Street_Address_Line1	Reci
Label	Recipient_State				
0	AE	3	3		3
	AK	101	101		101
	AL	1941	1932		1941
	AP	2	2		2
	AR	1014	1012		1014
	AZ	2153	2147		2153
	CA	10144	10098		10144
	CO	1173	1170		1173
	CT	1509	1502		1509

In [68]:

```
1 plt.figure(figsize = (25,10))
2 _ = sns.countplot(x="Recipient_State", data=df_final, hue='Label')
```



As we can see from the above chart, both classes have presence in each of the states. As such, we will preserve this feature and will not eliminate any rows.


In [69]:

```
1 #Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country
2 df_final.groupby(['Label', 'Applicable_Manufacturer_or_Applicable_GPO_Making_Pa
```

1	Australia	1	0
	Austria	540	0
	Barbados	16	0
	Belgium	29	4
	Canada	124	14
	Denmark	3681	14
	France	219	1
	Germany	3032	21
	Great Britain (UK)	817	37
	Ireland	135	0
	Israel	8	0
	Japan	100	7
	Netherlands	1242	51

In [70]:

```
1 plt.figure(figsize = (25,10))
2 _ = sns.countplot(x="Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Co
```



From the above table and chart, it is difficult to see the difference in category. For this, we will adopt a different approach where we will calculate the set difference to find out the categories that are missing in either of the classes.

In [71]:

```
1 #Countries not in Class 0
2 a = list(set(df_final[df_final.Label == 1].Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country))
3 a
```

Out[71]:

```
['Austria', 'Spain', 'New Zealand', 'Norway']
```

In [72]:

```
1 #Countries not in Class 1
2 b = list(set(df_final[df_final.Label == 0].Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country))
3 b
```

Out[72]:

```
['South Africa',
 'China',
 'Italy',
 'Mexico',
 'Hong Kong',
 'Korea (Republic of)',
 'Iceland']
```

In [73]:

```
1 #Removing rows with the above countries
2 df_final = df_final.loc[~df_final.Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country.isin(a+b)]
```



In [74]:

```
1 #Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State
2 df_final.groupby(['Label', 'Applicable_Manufacturer_or_Applicable_GPO_Making_Pa
```

NC	2681	389
NH	36	2
NJ	27990	976
NV	116	0
NY	10495	380
OH	105	41
OR	3039	86
PA	2122	57
PR	13	4
RI	63	33
TN	290	0
TX	911	102
...

In [75]:

```
1 plt.figure(figsize = (25,10))
2 _ = sns.countplot(x="Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_S
```



Similar to the previous feature, it is difficult to see the difference in category. For this, we will adopt a similar approach where we will calculate the set difference to find out the categories that are missing in either of the classes.

In [76]:

```
1 #Countries not in Class 0
2 a = list(set(df_final[df_final.Label == 1].Applicable_Manufacturer_or_Applicable
3 a
```

Out[76]:

['VT']

In [77]:

```
1 #Countries not in Class 1
2 b = list(set(df_final[df_final.Label == 0].Applicable_Manufacturer_or_Applicable_GPO_Making
3 b
```

Out[77]:

```
['KY', 'AR', 'IA', 'SD', 'NM', 'KS', 'LA', 'OK', 'AL', 'SC']
```

In [78]:

```
1 #Removing rows with the above countries
2 df_final = df_final.loc[~df_final.Applicable_Manufacturer_or_Applicable_GPO_Making
```

In [79]:


```
1 #Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1
2 df_final.groupby(['Label', 'Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1
```

Out[79]:

		Change_Type	Physician_Profile
Label	Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1		
0	Biological	9820	97
	Device	19255	190
	Drug	63185	63
	Medical Supply	84	
1	Biological	14362	4
	Device	11105	8
	Drug	44553	17
	Medical Supply	13	

In [80]:

```
1 plt.figure(figsize = (25,10))
2 _ = sns.countplot(x="Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1"
```



As we can see from the above chart, both classes have presence in each of the states. As such, we will

preserve this feature and will not eliminate any rows.

In [81]:

```
1 df_final.drop(columns=['Associated_Drug_or_Biological_NDC_1', 'Physician_Profile',
2                       'Applicable_Manufacturer_or_Applicable_GPO_Making_Payment'])
```

In [82]:

```
1 df_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 198880 entries, 0 to 199999
Data columns (total 20 columns):
Change_Type                                19
8880 non-null object
Recipient_City                             19
8759 non-null object
Recipient_State                            19
8679 non-null object
Recipient_Zip_Code                         19
8679 non-null object
Physician_Primary_Type                     19
7679 non-null object
Physician_Specialty                        19
7411 non-null object
Physician_License_State_code1              19
7676 non-null object
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name 19
8880 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name 19
8880 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State 18
8072 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country 19
8880 non-null object
Related_Product_Indicator                  19
8880 non-null object
Covered_or_Noncovered_Indicator_1          18
0469 non-null object
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1 16
2377 non-null object
Product_Category_or_Therapeutic_Area_1     15
7318 non-null object
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1 16
1606 non-null object
Total_Amount_of_Payment_USDollars          19
8880 non-null float64
Date_of_Payment                           19
8880 non-null object
Dispute_Status_for_Publication              19
8880 non-null object
Label                                       19
8880 non-null int64
dtypes: float64(1), int64(1), object(18)
memory usage: 36.9+ MB
```

In [83]:

```
1 df_final.head()
```

Out[83]:

	Change_Type	Recipient_City	Recipient_State	Recipient_Zip_Code	Physician_Primary_Type	Pr
0	UNCHANGED	MIDLAND	MI	48642-4600	Doctor of Osteopathy	
1	UNCHANGED	CARY	NC	27518	Doctor of Osteopathy	
2	UNCHANGED	Elgin	IL	60123	Medical Doctor	
3	UNCHANGED	PONCE	PR	00717	Medical Doctor	
4	UNCHANGED	TEMPE	AZ	85282-1892	Doctor of Osteopathy	

In [84]:

```
1 #Saving the final dataset
2 df_final.to_csv('data/df_final.csv')
```

Task 2

In the previous part we did intensive identification of the features to prevent information leak. In this section we will build a simple minimum viable model based on the initial features that were selected in previous part.

The final dataset of 200,000 rows that was created will be used here.

In [85]:

```
1 #Reading the Dataset created in the previous task
2 data = pd.read_csv('data/df_final.csv')
```

Since, building model on 200,000 is computationally expensive, we will further subsample the dataset to have 20,000 rows.

In [86]:

```
1 #Creating a sample with only 20% of all the values in the original data.
2 df_compressed = data.sample(frac= 0.1, random_state=42).reset_index(drop=True)
3 df_compressed.drop(columns='Unnamed: 0', inplace=True)
4 df_compressed.groupby(df_compressed['Label']).count()
```

Out[86]:

	Change_Type	Recipient_City	Recipient_State	Recipient_Zip_Code	Physician_Primary_Type
Label					
0	9963	9963	9963	9963	9916
1	9925	9911	9899	9899	9854

As you can see from the above, we have created a balanced dataset through undersampling. Both classes - Class 1 and Class 2 have almost same number of rows in the dataset.

Now, we will preprocess the Date_of_time variable. Since the data is from 2017, we will remove the year and the day and just keep the month, which will be treated as a categorical variable.

In [87]:

```
1 #Converting the Date_of_Payment to datetime
2 df_compressed.Date_of_Payment = pd.to_datetime(df_compressed.Date_of_Payment)
3 #Adding the Month of Payment
4 df_compressed['Payment_Month'] = df_compressed.Date_of_Payment.dt.month
5 #Dropping Date_of_Payment (Year is 2017)
6 df_compressed.drop(columns='Date_of_Payment', inplace=True)
```

In [88]:

```
1 df_compressed.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19888 entries, 0 to 19887
Data columns (total 20 columns):
Change_Type                                19
888 non-null object
Recipient_City                             19
874 non-null object
Recipient_State                             19
862 non-null object
Recipient_Zip_Code                         19
862 non-null object
Physician_Primary_Type                     19
770 non-null object
Physician_Specialty                        19
743 non-null object
Physician_License_State_code1              19
770 non-null object
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name 19
888 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name 19
888 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State 18
867 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country 19
888 non-null object
Related_Product_Indicator                  19
888 non-null object
Covered_or_Noncovered_Indicator_1          18
085 non-null object
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1 16
284 non-null object
Product_Category_or_Therapeutic_Area_1     15
770 non-null object
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1 16
216 non-null object
Total_Amount_of_Payment_USDollars          19
888 non-null float64
Dispute_Status_for_Publication              19
888 non-null object
Label                                       19
888 non-null int64
Payment_Month                             19
888 non-null int64
dtypes: float64(1), int64(2), object(17)
memory usage: 3.0+ MB
```

In [89]:

```
1 #Converting Payment_Month to object as it is a categorical variable
2 df_compressed['Payment_Month'] = df_compressed['Payment_Month'].astype(object)
```

In [90]:

```
1 x=df_compressed.drop(columns=['Label'])
2 y=df_compressed.Label
```

We will define the categorical and continuous features. In our dataset, there is only one continuous feature - Total_Amount_of_Payment_USDollars. Rest all the features are categorical.

In [91]:

```
1 #Categorical Features
2 cat_features = list(X.columns)
3 cat_features.remove('Total_Amount_of_Payment_USDollars')
4
5 #Continuous Features
6 cont_features = ['Total_Amount_of_Payment_USDollars']
```

The NaN values for the categorical features will be encoded into a new category 'NA'. There are no null values for the continuous feature.

In [92]:

```
1 X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19888 entries, 0 to 19887
Data columns (total 19 columns):
Change_Type                                          19
888 non-null object
Recipient_City                                      19
874 non-null object
Recipient_State                                     19
862 non-null object
Recipient_Zip_Code                                  19
862 non-null object
Physician_Primary_Type                             19
770 non-null object
Physician_Specialty                                19
743 non-null object
Physician_License_State_code1                      19
770 non-null object
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name 19
888 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name 19
888 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State 18
867 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country 19
888 non-null object
Related_Product_Indicator                          19
888 non-null object
Covered_or_Noncovered_Indicator_1                  18
085 non-null object
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1 16
284 non-null object
Product_Category_or_Therapeutic_Area_1             15
770 non-null object
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1 16
216 non-null object
Total_Amount_of_Payment_USDollars                  19
888 non-null float64
Dispute_Status_for_Publication                     19
888 non-null object
Payment_Month                                       19
888 non-null object
dtypes: float64(1), object(18)
memory usage: 2.9+ MB
```

In [93]:

```
1 # Splitting the data in training set and test set
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

Proceeding to building the baseline model - Logistic Regression.

- We will use limited featured while building the model. As such only use a subset of columns which includes randomly selected 8 features. Columns that would be included are 'Change_Type', 'Recipient_Zip_Code', 'Physician_Specialty',


```
'Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country', 'Related_Product_Indicator',
'Total_Amount_of_Payment_USDollars', 'Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1',
'Covered_or_Noncovered_Indicator_1'
```

- Note that above features are chosen on a random basis as this is a baseline model.
- In terms of preprocessing, a new category would be created to encode the null values in each of the column.
- Since this is the baseline model, we will use the default parameters specified in scikit learn for Logistic Regression.

The NaN values for the categorical features will be encoded into a new category 'NA'. There are no null values for the continuous feature.

In [101]:

```
# Replacing the NaN values to NA which is a new category created
def null_categories(Xa):
    3 Xa = Xa[['Change_Type', 'Recipient_Zip_Code', 'Physician_Specialty', 'Applicable_
    4           'Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1', 'Covered_or
    5 return Xa.replace({np.nan: 'NA'})
```

In [102]:

```
1 # Build a pipeline for dealing with categorical variables and continuous variables
2 from sklearn.compose import make_column_transformer
3 from sklearn.compose import ColumnTransformer
4 from sklearn.preprocessing import FunctionTransformer
5 from sklearn.preprocessing import OneHotEncoder
6
7 from sklearn.pipeline import make_pipeline
8
9 preprocess_baseline_nan = ColumnTransformer(
10     [('func', FunctionTransformer(func=null_categories, validate=False), cat_fea
11
12 preprocess_baseline_ohe = ColumnTransformer([('ohe', OneHotEncoder(handle_unknown=
13                                             np.arange(0,8))], remainder='passthru')
```

In [103]:

```
1 #Model Building - Baseline Model (Logistic Regression)
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.model_selection import cross_val_score
4
5 pipeline_baseline = make_pipeline(preprocess_baseline_nan, preprocess_baseline_ohe)
6 logreg_scores = cross_val_score(pipeline_baseline, X_train, y_train, cv=5)
7 print('Logistic Regression - Accuracy (Baseline Model): ', np.mean(logreg_scores))
```

Logistic Regression - Accuracy (Baseline Model): 0.8429211620271678

In [104]:

```
1 precision_baseline = cross_val_score(pipeline_baseline, X_train, y_train, cv=5)
2 recall_baseline = cross_val_score(pipeline_baseline, X_train, y_train, cv=5, scoring='recall')
3 print('Logistic Regression - Precision (Baseline Model): ', np.mean(precision_baseline))
4 print('Logistic Regression - Recall (Baseline Model): ', np.mean(recall_baseline))
```

Logistic Regression - Precision (Baseline Model): 0.873843083096775

Logistic Regression - Recall (Baseline Model): 0.8009128772287095

As we can see from the above analysis, a baseline linear model that selects 8 random features achieves an accuracy of 86%. In the subsequent parts, we will try to beat this accuracy using better feature selection mechanisms.

Task 3

In this section, we will create derived features, do more in-depth processing and data cleaning. This will particularly revolve around encoding the categorical variables in different ways. Using this we will check whether this improves the model.

We will follow the following approaches to preprocessing:

- All null values will be encoding in a different category, similar to Task 2.
- Target Encoding will be used to encode categorical features.
- Standardize the continuous feature 'Total_Amount_of_Payment_USDollars'.
- All the features would be used in model building.

In [105]:

```
1 X.nunique().sort_values()
```

Out[105]:

```
Covered_or_Noncovered_Indicator_1
2
Related_Product_Indicator
2
Dispute_Status_for_Publication
2
Change_Type
3
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1
4
Physician_Primary_Type
5
Payment_Month
12
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country
14
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State
32
Recipient_State
52
Physician_License_State_code1
52
Physician_Specialty
90
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name
5
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name
5
Product_Category_or_Therapeutic_Area_1
6
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1
17
Recipient_City
34
Recipient_Zip_Code
87
Total_Amount_of_Payment_USDollars
94
dtype: int64
```

In this section, we will drop the obvious correlation features:

- Recipient_City & Recipient_Zip_Code: Correlated with Recipient_State
- Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Country: Correlated with Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State

In [106]:

```
1 #Dropping redundant features which are obviously correlated
2 X.drop(columns=['Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Count:',
3                'Recipient_Zip_Code', 'Recipient_City'], inplace=True)
4
5 #Finding the Categorical Features
6 cat_features = list(X.columns)
7 cat_features.remove('Total_Amount_of_Payment_USDollars')
```


In [107]:

```
1 #Defining a function to replace the null values with NA
2 def null_categories_q3(Xa):
3     return Xa.replace({np.nan: 'NA'})
```

In [110]:

```
1 from sklearn.preprocessing import StandardScaler
2 from category_encoders.target_encoder import *
3 #Checking the Correlation of other variables
4
5 #Creating a ColumnTransformer
6 preprocess_nan = ColumnTransformer(
7     [('func', FunctionTransformer(func=null_categories_q3, validate=False), cat_
8     ('scaler', StandardScaler(), cont_features)], remainder='passthrough')
9
10 #Transforming through the pipeline
11 pipeline_q3 = make_pipeline(preprocess_nan, TargetEncoder(impute_missing=False))
12 pipeline_q3.fit_transform(X, y.values)
13
14 #Scaling the Target Encoded Variables
15 from sklearn.preprocessing import scale
16 X_transformed = pipeline_q3.fit_transform(X, y.values)
17 X_scaled = scale(X_transformed)
18 cov = np.cov(X_scaled, rowvar=False)
19
20 #Plotting the Covariance Plot
21 plt.figure(figsize=(8, 8), dpi=100)
22 plt.imshow(cov)
23 plt.xticks(range(X.shape[1]), cat_features+cont_features, rotation=90)
24 plt.yticks(range(X.shape[1]), cat_features+cont_features)
25 plt.title('Covariance Between Variables');
```

/Users/Sarang/anaconda3/envs/coms007/lib/python3.6/site-packages/ipykernel_launcher.py:17: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by the scale function.



As we can see from above, there are some evident correlations:

- 'Physician_License_State_code1' is highly correlated with 'Recipient_State'
- 'Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name' is highly correlated with 'Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name'
- 'Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1' is highly correlated with 'Product_Category_or_Therapeutic_Area_1'

To tackle this, we will remove the correlated feature that has less number of null values. In case of a tie, we will remove the features that has more unique values to better generalize the data.

In [111]:

```
1 X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19888 entries, 0 to 19887
Data columns (total 16 columns):
Change_Type                                1988
8 non-null object
Recipient_State                            1986
2 non-null object
Physician_Primary_Type                    1977
0 non-null object
Physician_Specialty                       1974
3 non-null object
Physician_License_State_code1             1977
0 non-null object
Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name 1988
8 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name 1988
8 non-null object
Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State 1886
7 non-null object
Related_Product_Indicator                 1988
8 non-null object
Covered_or_Noncovered_Indicator_1         1808
5 non-null object
Indicate_Drug_or_Biological_or_Device_or_Medical_Supply_1 1628
4 non-null object
Product_Category_or_Therapeutic_Area_1    1577
0 non-null object
Name_of_Drug_or_Biological_or_Device_or_Medical_Supply_1 1621
6 non-null object
Total_Amount_of_Payment_USDollars         1988
8 non-null float64
Dispute_Status_for_Publication            1988
8 non-null object
Payment_Month                             1988
8 non-null object
dtypes: float64(1), object(15)
memory usage: 2.4+ MB
```

In [112]:

```
1 X.drop(columns=['Physician_License_State_code1',
2               'Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_Name'],
3
4 #Updating the Categorical Features
5 cat_features = list(X.columns)
6 cat_features.remove('Total_Amount_of_Payment_USDollars')
```

In [113]:

```
1 # Splitting the data in training set and test set
2 from sklearn.model_selection import train_test_split
3 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

In [114]:

```
1 #Creating a column transformer with null_categories_q3 function
2 preprocess_nan = ColumnTransformer(
3     [('func', FunctionTransformer(func=null_categories_q3, validate=False), cat
4     ('scaler', StandardScaler(), cont_features)], remainder='passthrough')
```

In [115]:

```
1 #Creating Pipeline to standardize continuous features and target encode categori
2 pipeline_q3 = make_pipeline(preprocess_nan, TargetEncoder(impute_missing=False))
3 X_te = pipeline_q3.fit_transform(X_train, y_train.values)
4
5 #Running a 5-fold cross validation on the baseline model - Logistic Regression.
6 logreg_scores = cross_val_score(LogisticRegression(), X_te, y_train, cv=5)
7 print('Logistic Regression - Accuracy (Target Based): ', np.mean(logreg_scores))
```

Logistic Regression - Accuracy (Target Based): 0.8913911803460562

In [116]:

```
1 #Calculating the precision and recall
2 precision_q3 = cross_val_score(LogisticRegression(), X_te, y_train, cv=5, scor
3 recall_q3 = cross_val_score(LogisticRegression(), X_te, y_train, cv=5, scoring=
4 print('Logistic Regression - Precision (Target Based): ', np.mean(precision_q3)
5 print('Logistic Regression - Recall (Target Based): ', np.mean(recall_q3))
```

Logistic Regression - Precision (Target Based): 0.897625720911762

Logistic Regression - Recall (Target Based): 0.8831265300374792

It looks like the Accuracy, Precision and Recall improve significantly. Let us now one-hot-encode the categorical features instead of target encoding and see how the results are affected.

In [117]:

```
1 #Creating Column Transformer for one-hot-encoding
2 preprocess_nan = ColumnTransformer(
3     [('func', FunctionTransformer(func=null_categories_q3, validate=False), cat
4     ('scaler', StandardScaler(), cont_features)], remainder='passthrough')
5
6 preprocess_baseline_ohe = ColumnTransformer([('ohe', OneHotEncoder(handle_unkno
7     np.arange(0, len(X_train.columns)-1))
```

In [118]:

```
1 #Creating pipeline for one-hot-encoding
2 pipeline_q3_1 = make_pipeline(preprocess_nan, preprocess_baseline_ohe, Logistic
3 logreg_scores = cross_val_score(pipeline_q3_1, X_train, y_train, cv=5)
4 print('Logistic Regression - Accuracy (OHE): ', np.mean(logreg_scores))
```

Logistic Regression - Accuracy (OHE): 0.9027888331059764

In [119]:

```
1 #Calculating the precision and recall
2 precision_q3_1 = cross_val_score(pipeline_q3_1, X_train, y_train, cv=5, scoring='precision')
3 recall_q3_1 = cross_val_score(pipeline_q3_1, X_train, y_train, cv=5, scoring='recall')
4 print('Logistic Regression - Precision (OHE): ', np.mean(precision_q3_1))
5 print('Logistic Regression - Recall (OHE): ', np.mean(recall_q3_1))
```

Logistic Regression - Precision (OHE): 0.9145777401938282

Logistic Regression - Recall (OHE): 0.8882320710298461

It appears that there isn't much difference in the way features are encoded. Both Target Based Encoding and One Hot Encoding give almost the same results. However, since there are a lot of categories, we will go ahead with using target encoding in the following.

Task 4

In this section, we will implement two classification models - Random Forest and Gradient Boosting. The preprocessing and feature engineering will be done as per the requirement of the model.

The dataset that would be used is the X_train and y_train that was created after we removed the features in the above tasks.

The following preprocessing and feature engineering steps would be done for both RandomForest and GradientBoosting. The steps are same as both of them are ensemble tree based models:

- Create a new category for the null values.
- Target encode the categorical variables.
- We will not scale the data as both are tree based model and scaling can be computationally expensive.

In [120]:

```
1 #Defining a function to replace the null values with NA
2 def null_categories_q4(Xa):
3     return Xa.replace({np.nan: 'NA'})
4
5 #Creating a column transformer with null_categories_q3 function
6 preprocess_nan_q4 = ColumnTransformer(
7     [('func', FunctionTransformer(func=null_categories_q4, validate=False), cat_1
```

First, we will implement the **Random Forest Classifier**. The parameters that we would be tuning are - n_estimators and max_depth over 3 values each.

In [121]:

```
1 from sklearn.model_selection import GridSearchCV
2 from sklearn.pipeline import Pipeline
3 from sklearn.ensemble import RandomForestClassifier
4
5 scoring = {'auc': 'roc_auc', 'accuracy': 'accuracy', 'average_precision': 'average_precision',
6           'precision': 'precision', 'recall': 'recall'}
7
8 #Creating Pipeline for Random Forest
9 pipeline_rf = Pipeline(steps=[('nan', preprocess_nan_q4), ('te', TargetEncoder()),
10                             ('rf', RandomForestClassifier())])
11
12 #Defining the parameters over which grid-search would run: Note only two parameters
13 #the model is computationally expensive - it makes the kernel crash on our machine
14
15 parameters = {'rf__n_estimators': [100, 500, 1000], 'rf__max_depth': [None, 3, 7]}
16
17 #Running Grid Search
18 gcv_rf = GridSearchCV(pipeline_rf, param_grid=parameters, cv=3, scoring=scoring)
19 gcv_rf.fit(X_train, y_train.values)
20
21 print("RF best parameters: {}".format(gcv_rf.best_params_))
```

RF best parameters: {'rf__max_depth': None, 'rf__n_estimators': 500}

In [122]:

```
1 print("accuracy RF: {}".format(np.mean(gcv_rf.cv_results_['mean_test_accuracy'])))
2 print("average_precision RF: {}".format(np.mean(gcv_rf.cv_results_['mean_test_average_precision'])))
3 print("precision RF: {}".format(np.mean(gcv_rf.cv_results_['mean_test_precision'])))
4 print("recall RF: {}".format(np.mean(gcv_rf.cv_results_['mean_test_recall'])))
```

accuracy RF: 0.9403921218080511
average_precision RF: 0.979491694038916
precision RF: 0.9270312094666393
recall RF: 0.9559226800726699

Now, we will implement a **Gradient Boosting Classifier**. We would be tuning learning rate over 3 values (note that fewer hyperparameters are tuned due to computational limitations).

In [123]:

```
1 from sklearn.ensemble import GradientBoostingClassifier
2
3 #Creating Pipeline for Gradient Boosting
4 pipeline_gb = Pipeline(steps=[('nan', preprocess_nan_q4), ('te', TargetEncoder()),
5                             ('gb', GradientBoostingClassifier())])
6
7 #Defining the parameters over which grid-search would run
8 parameters = {'gb__learning_rate': [.2, .1, .05]}
9
10 #Running Grid Search
11 gcv_gb = GridSearchCV(pipeline_gb, param_grid=parameters, cv=3, scoring=scoring)
12 gcv_gb.fit(X_train, y_train.values)
13
14 print("GB best parameters: {}".format(gcv_gb.best_params_))
```

GB best parameters: {'gb__learning_rate': 0.2}

In [124]:

```
1 print("accuracy GB: {}".format(np.mean(gcv_gb.cv_results_['mean_test_accuracy']
2 print("average_precision GB: {}".format(np.mean(gcv_gb.cv_results_['mean_test_a
3 print("precision GB: {}".format(np.mean(gcv_gb.cv_results_['mean_test_precision
4 print("recall GB: {}".format(np.mean(gcv_gb.cv_results_['mean_test_recall'])))
```

```
accuracy GB: 0.9434835076427998
average_precision GB: 0.9824335714538619
precision GB: 0.9321416412784284
recall GB: 0.9563851645505409
```

As we can see from the above results, both Random Forest and Gradient Boosting perform equally well. Note that the results that we get in this part are higher than what we got in Task 2 (Baseline) and improved Task 3 (Improved Baseline).

Task 5

In this section, we will identify features that are important to our best model. Since there is not much difference between Gradient Boosting and Random Forest as per Task 4, we will go ahead and choose Gradient Boosting to be the model that we will use in this task with the best parameter - learning_rate = 0.2

We will also study which features are most influential, and which features could be removed without decrease in performance.

In [125]:














```
1 preprocess_nan_q5 = ColumnTransformer(
2     [('func', FunctionTransformer(func=null_categories_q4, validate=False), cat_
```

First, we will plot the **Partial Dependence Curves** to see the dependence of each of the variable on the label individually.

In [155]:

```
1 from sklearn.ensemble.partial_dependence import plot_partial_dependence
2 x = preprocess_nan_q5.fit_transform(X_train)
3
4 _ = plt.figure(figsize=(100,100))
5
6 fig, axs = plot_partial_dependence(
7     gcv_gb.best_estimator_.named_steps['gb'], TargetEncoder().fit_transform(x, y),
8     np.argsort(gcv_gb.best_estimator_.named_steps['gb'].feature_importances_)[-
9     feature_names=cat_features+cont_features, n_jobs=-1)
10
11 plt.subplots_adjust(top=5, left=0, right=2)
```

<Figure size 7200x7200 with 0 Axes>



The curves for 'Dispute_Status_for_Publication', 'Change_Type', 'Physician_Primary_Type', 'Related_Product_Indicator' and 'Covered_or_Noncovered_Indicator_1' are almost flat. This suggests that these features do not individually impact the target.

Now let us use **Recursive Feature Elimination**. This iterative model-based selection would help us drop features in a recursive fashion while taking into account the impact of other features.

In [160]:


```
1 from sklearn.feature_selection import RFE
2
3 # create ranking among all features by selecting only one
4 rfe = RFE(gcv_gb.best_estimator_.named_steps['gb'], n_features_to_select=1)
5 rfe.fit(TargetEncoder().fit_transform(x, y_train), y_train)
6 rfe.ranking_
```

Out[160]:

```
array([ 9, 12,  7,  8,  2,  3, 10,  4,  5,  6, 13, 11,  1])
```

In [173]:

```
1 #Plotting the importance of features obtained in RFE
2 fig = plt.figure(figsize=(12,6))
3 _ = plt.plot(1-rfe.ranking_, 'o', c='r')
4 _ = plt.ylabel("1-RFE Rankings")
5 _ = plt.xticks(range(X.shape[1]), cat_features+cont_features, rotation=90)
6 _ = plt.title('Feature Importance through RFE')
```



The results obtained in Partial Dependence Curve and Recursive Feature Elimination are slightly different. This might be due to the fact that Partial Dependence Curve ignore all the other features while reporting the importance whereas RFE adopts a iterative approach which takes into account the effect of other features.

Next, we will employ **Recursive Feature Elimination Cross Validation (RFECV)** that allows you to us efficient grid search for the number of features to keep.

In [189]:

```
1 #Performing Recursive Feature Elimination Cross Validation
2 from sklearn.feature_selection import RFECV
3
4 X_train_transformed = TargetEncoder().fit_transform(x, y_train)
5
6 #Running RFECV to get the best parameters
7 rfe = RFECV(gcv_gb.best_estimator_.named_steps['gb'], cv=5)
8 rfe.fit(X_train_transformed, y_train)
9 print(rfe.support_)
10 print(list(np.array(cat_features+cont_features)[rfe.support_]))
```

```
[ True False  True  True  True  True  True  True  True  True False  Tr
ue
 True]
['Change_Type', 'Physician_Primary_Type', 'Physician_Specialty', 'Subm
itting_Applicable_Manufacturer_or_Applicable_GPO_Name', 'Applicable_Ma
nufacturer_or_Applicable_GPO_Making_Payment_State', 'Related_Product_I
ndicator', 'Covered_or_Noncovered_Indicator_1', 'Indicate_Drug_or_Biol
ogical_or_Device_or_Medical_Supply_1', 'Product_Category_or_Therapeuti
c_Area_1', 'Payment_Month', 'Total_Amount_of_Payment_USDollars']
```

As per RFECV, the above features are useful for the prediction. Now, we will build a model based on these features to see if the model performance improves.

In [192]:

```
#Building Gradient Boosting Model based on the features selected by RFECV
from sklearn.model_selection import cross_validate
3
pipe_rfe_gb = make_pipeline(RFECV(gcv_gb.best_estimator_.named_steps['gb'], cv=5), gcv_gb)
pipe_rfe_gb_scores = cross_validate(pipe_rfe_gb, X_train_transformed, y_train, cv=5,
```

In [195]:

```
1 #Calculating the results
2 print("accuracy GB-RFECV: {}".format(np.mean(pipe_rfe_gb_scores['test_accuracy'])))
3 print("average_precision GB-RFECV: {}".format(np.mean(pipe_rfe_gb_scores['test_avg_precision'])))
4 print("precision GB-RFECV: {}".format(np.mean(pipe_rfe_gb_scores['test_precision'])))
5 print("recall GB-RFECV: {}".format(np.mean(pipe_rfe_gb_scores['test_recall'])))
6 print("auc GB-RFECV: {}".format(np.mean(pipe_rfe_gb_scores['test_auc'])))
```

```
accuracy GB-RFECV: 0.9410030811328248
average_precision GB-RFECV: 0.9806453397544246
precision GB-RFECV: 0.927698306795922
recall GB-RFECV: 0.9563411252410148
auc GB-RFECV: 0.9833411665443841
```

As we can see from the above, there is not much difference from the Gradient Boosting Model that was created in Task 4. This is probably due to the fact that the performance of the model in Task 4 was already very good. Hence, there might be a ceiling effect limiting how much better scores we can obtain.

Task 6

In this task, we will create an 'explainable' model that will try to achieve the same performance results as our

best model - Gradient Boosting built in the previous sections.

For this task, we will develop a simple tree. We will perform grid search to tune our parameters, use feature selection to reduce the number of features and also perform preprocessing suitable for Decision Tree. Since this model will be an explainable one, it will have small number of leaves/less depth.

We will focus on explaining the model globally i.e. we will identify the features that are important to the model.

In [238]:

```
from sklearn.tree import DecisionTreeClassifier, export_graphviz
pipeline_tree = Pipeline(steps=[('nan', preprocess_nan_q4), ('te', TargetEncoder())],

#Defining the parameters over which grid-search would run: Note only two parameters
#the model is computationally expensive - it makes the kernel crash on our machines

parameters = {'tree__max_depth': [None, 4, 10]}

#Running Grid Search
gcv_tree = GridSearchCV(pipeline_tree, param_grid= parameters, cv=3, scoring=scoring)
gcv_tree.fit(X_train, y_train.values)

print("Tree best parameters: {}".format(gcv_tree.best_params_))
```

Tree best parameters: {'tree__max_depth': 4}

In [360]:

```
1 print("accuracy tree: {}".format(np.mean(gcv_tree.cv_results_['mean_test_accuracy'])))
2 print("average_precision tree: {}".format(np.mean(gcv_tree.cv_results_['mean_test_average_precision'])))
3 print("precision tree: {}".format(np.mean(gcv_tree.cv_results_['mean_test_precision'])))
4 print("recall tree: {}".format(np.mean(gcv_tree.cv_results_['mean_test_recall'])))
5 print("auc tree: {}".format(np.mean(gcv_tree.cv_results_['mean_test_auc'])))
```

accuracy tree: 0.9332037186019487
average_precision tree: 0.9214610912652677
precision tree: 0.9245028487325908
recall tree: 0.9433996366227194
auc tree: 0.9468890628021424

Note from the above that the performance of a single decision tree is nearly as good as best Gradient Boosting Model.

Now let us visualize the decision tree with the best parameters.


In [275]:

```

1 #Visualizing the Decision Tree:
2 #source: https://medium.com/@rnbrown/creating-and-visualizing-decision-trees-wi
3
4 from sklearn.externals.six import StringIO
5 from IPython.display import Image
6 from sklearn.tree import export_graphviz
7 import pydot
8 import graphviz
9 import pydotplus
10 dot_data = StringIO()
11 export_graphviz(gcv_tree.best_estimator_.named_steps['tree'], out_file=dot_data,
12                 filled=True, rounded=True,
13                 special_characters=True, feature_names=(cat_features+cont_features),
14                 graph = pydotplus.graph_from_dot_data(dot_data.getvalue()))
15 Image(graph.create_png(),width=2000, height=2000)

```

Out[275]:



As we can see from the above tree, the model is very explainable. We can see that the tree is split on different features alongside the Gini coefficient allowing us to see which coefficient leads to maximum reduction in the Gini coefficient.

PS: Please ignore the small size of the chart due to large feature names, affecting the aspect ratio


Now, let us plot the Feature Importances

In [292]:

```

1 #Plotting Feature Importances.
2 _ = plt.figure(figsize=(12,6))
3 _ = plt.barh(range(13), gcv_tree.best_estimator_.named_steps['tree'].feature_importances_)
4 _ = plt.yticks(range(13), (cat_features+cont_features))
5 _ = plt.title('Feature Importance - Decision Tree')
6 _ = plt.xlabel('Importance')
7 _ = plt.ylabel('Feature')

```



Now we will build even a simpler model, where we will take only the 4 features as per the above chart that are

important. We will eliminate all the other features:

- 'Total_Amount_of_Payment_USDollars'
- 'Product_Category_or_Therapeutic_Area_1'
- 'Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State'
- 'Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name'

In [370]:

```
1 #Creating reduce training set
2
3 X_train_reduced = X_train[['Total_Amount_of_Payment_USDollars', 'Product_Category_or_Therapeutic_Area_1',
4 'Applicable_Manufacturer_or_Applicable_GPO_Making_Payment_State', 'Submitting_Applicable_Manufacturer_or_Applicable_GPO_Name']]
5
6 X_train_reduced.reset_index(inplace=True)
7 X_train_reduced.drop(columns='index', inplace=True)
```

In [372]:

```
preprocess_nan_q6 = ColumnTransformer(
    [ ('func', FunctionTransformer(func=null_categories_q4, validate=False), X_train_reduced)],
    remainder='passthrough')

pipeline_tree_update = Pipeline(steps=[ ('nan', preprocess_nan_q6), ('te', TargetEncoder()) ])

#Defining the parameters over which grid-search would run: Note only two parameters
#the model is computationally expensive - it makes the kernel crash on our machines

parameters = {'tree__max_depth': [2, 3, 4]}

# # #Running Grid Search
gcv_tree_updated = GridSearchCV(pipeline_tree_update, param_grid= parameters, cv=3,
gcv_tree_updated.fit(X_train_reduced, y_train.values)

print("Tree best parameters: {}".format(gcv_tree_updated.best_params_))
```

Tree best parameters: {'tree__max_depth': 4}

In [373]:

```
1 print("accuracy tree (4 features): {}".format(np.mean(gcv_tree_updated.cv_results_['accuracy'])))
2 print("average_precision tree (4 features): {}".format(np.mean(gcv_tree_updated.cv_results_['average_precision'])))
3 print("precision tree (4 features): {}".format(np.mean(gcv_tree_updated.cv_results_['precision'])))
4 print("recall tree (4 features): {}".format(np.mean(gcv_tree_updated.cv_results_['recall'])))
5 print("auc tree (4 features): {}".format(np.mean(gcv_tree_updated.cv_results_['auc'])))
```

```
accuracy tree (4 features): 0.9092026459283097
average_precision tree (4 features): 0.8957664351150787
precision tree (4 features): 0.8779306427984803
recall tree (4 features): 0.9515940559075123
auc tree (4 features): 0.9352025444704372
```

As we can see from the above, using only the 4 features does not decrease the performance to a huge extent. But at the same time we get a very good explainable model.

Let us now visualize this updated decision tree.

In [383]:

```
1 #Visualizing the Decision Tree:
2 #source: https://medium.com/@rnbrown/creating-and-visualizing-decision-trees-wi
3
4 dot_data = StringIO()
5 export_graphviz(gcv_tree_updated.best_estimator_.named_steps['tree'], out_file=dot_data,
6                 filled=True, rounded=True,
7                 special_characters=True, feature_names=X_train_reduced.columns)
8 graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
9 Image(graph.create_png(),width=2000, height=2000)
```

Out[383]:

